

Large-scale hybrid ad hoc network for mobile platforms: Challenges and Experiences

Nirmit Desai, Wendy Chong, Heather Achilles, Shahrokh Daijavad
IBM T. J. Watson Research Center
Yorktown Heights, NY, USA
{nirmit.desai, wendych, hachilles, shahrokh}@us.ibm.com

Thomas La Porta
Dept. of Computer Science and Engineering
Penn State University
University Park, PA, USA
t1p@cse.psu.edu

Abstract—Peer-to-peer (p2p) networks and Mobile ad hoc networks (MANET) have been widely studied. However, a real-world deployment for the masses has remained elusive. Ever-increasing density of mobile devices, especially in urban areas, has given rise to new applications of p2p communication. However, the modern smartphone platforms have limited support for such communications. Further, the issues of battery life, range, and trust remain unaddressed. A key question then is, what kinds of applications can the modern mobile platforms support and what challenges remain? This paper identifies a class of applications and presents a novel center-to-peer-to-peer (c2p2p) architecture called Mesh Network Alerts (MNA) to support them. We describe our experiences in deploying MNA as a real-world system to millions of users for relaying severe weather information along with the challenges faced, and the approaches for addressing them.

Index Terms—peer-to-peer systems, mobile ad hoc network, delay-tolerant network

I. INTRODUCTION

Mobile devices with programmable platforms such as android and iOS have steadily grown over the last decade, surpassing the 2 billion mark¹. MANETs have been widely studied given their decentralized nature and potential for new applications [11], [18]. Most of the prior work on p2p networks has focused on analytical and simulation-based study of MANET behavior [12], [13], [21]. Given the outstanding practical challenges in deploying a large number of physical nodes, real-world implementations have been limited and have not reached mass scale [10]. However, with the growth of smartphones, large-scale real-world implementations may become feasible. This paper describes a real-world implementation of a p2p delay-tolerant network, called Mesh Network Alerts (MNA), for relaying severe weather information to millions of mobile device users as part of the Weather Channel app [8] on both android and iOS platforms.

Before describing MNA, it is critical to identify applications that need p2p communication, given pervasive Internet connectivity. Doing so enables us to define key characteristics of such applications and focus on the challenges in meeting them. This paper focuses on two separate classes of applications: communication in (a) disaster-affected or remote areas and (b) congested networks in densely populated areas, e.g., sports

arenas. The following are the key characteristics in these scenarios:

- CH₀. No communication infrastructure such as WiFi access points to fall back on
- CH₁. Network nodes are mobile, pattern of mobility is not predictable
- CH₂. New information may arrive at any time
- CH₃. Trustworthy information is scarce, misinformation and rumours are common place
- CH₄. Small payloads suffice in many cases and information retains value for a few minutes
- CH₅. Device battery is a scarce resource, power supply for recharging may not be available
- CH₆. Devices are owned by citizens, deployment of special-purpose devices is cost prohibitive

The above needs are well-recognized in the industry as well as academia with several ambitious attempts to address them, e.g., Google Loon project² and Facebook Aquila³, though with limited impact. Leveraging user mobile devices as peer nodes for a large-scale deployment has been another theme in the prior works, e.g., the Serval project [6]. Serval mesh enables p2p communication over on-device WiFi radio, but requires root access to the device via jail-breaking. Although significant lessons have been learned through these attempts, a mass-scale p2p network for such applications remains elusive.

A vast majority of the literature has focused on a traditional model of stateful, fully decentralized, reliable networking. Specifically, the nodes maintain connectivity with peers and routing is optimized with techniques based on link state or distance vectors [5], [17] focusing on optimizing the network utilization.

Given the application characteristics above, this paper identifies main practical challenges associated with modern device platforms and finds novel ways to overcome them. This leads to MNA — a new paradigm in p2p networking that employs a hybrid, delay-tolerant, and zero-routing overhead architecture. Unlike previous MANET architectures, MNA is a hybrid of a centralized and a decentralized architecture, called *center-to-peer-to-peer (c2p2p)*. A central service is leveraged as the trusted source of information while the information

¹<https://www.statista.com/statistics/330695/number-of-smartphone-users-worldwide/>

²<https://loon.co>

³https://en.wikipedia.org/wiki/Facebook_Aquila

is propagated in a decentralized fashion by peers. Such an architecture allows MNA to bypass the key issue of trust deficit in open decentralized systems and scale while retaining the ability of infrastructure-less communication in many application scenarios.

MNA is implemented on both Android and iOS as an SDK and integrated with the Weather Channel mobile apps. With extensive experiments and experience of deploying to almost 10 million users, MNA represents a way forward for large-scale p2p networks.

This paper makes the following key contributions:

- Identification of a class of applications for p2p and their key characteristics
- A deeper investigation of the practical challenges in supporting the above class of applications
- A novel c2p2p architecture implementation using multiple radio channels for modern mobile platforms (Android and iOS)
- Experimental evaluation and deployment statistics

In the following, Section II describes the practical challenges. Section III outlines the architectural details of MNA along with platform-specific implementation issues for Android and iOS in addressing the challenges. Experimental evaluation and deployment statistics are presented in Section V. A deeper look at the literature and contrast to MNA is summarized in Section VI with conclusions in Section VII.

II. CHALLENGES

We present the main challenges for modern mobile device platforms in supporting the classes of applications described above. These have been uncovered via extensive experiments and in some cases include direct feedback from the developers of Android and iOS.

A. Operating system restrictions

Due to CH₂, even when a user is not interacting with an app, or worse yet, when the device is not being used at all, the devices must continue to discover peer devices to receive and forward information. Although modern mobile operating systems such as Android and iOS offer APIs to discover and advertise information to peer devices over WiFi and Bluetooth interfaces, peer-to-peer connections do not work reliably when the same APIs are accessed while the app is in the background. Further, on Android, each WiFi p2p connection must be explicitly approved by the device user. Prior works widely document these challenges and take the approach of having special access on the devices, e.g., jail-breaking or rooting [6]. Clearly, such an approach does not scale to mass adoption.

B. Power constraints

Since devices may be offline when new weather information arrives, MNA on each peer must remain active at all times to be able to discover new information as soon as it arrives. Further, as there is no back up infrastructure (CH₀) and a set

of peers in range can change at any time (CH₁), each peer is responsible for constantly forwarding available information to other peers via advertisements. However, due to CH₅, the MNA activity must keep the device battery consumption to a minimum. This is a challenge because advertising and discovery are power-hungry operations over the radio channels.

C. Testing p2p networks

Given the heterogeneity of devices owned by users and operating system distributions (CH₆), it is challenging to test whether or not MNA functions as expected on a single device. Further, running test scenarios on a p2p network at large-scale is non-trivial given that a large number of devices need to take specific coordinated action followed by coordinated observations to determine whether a test passes or fails. Further, since range and mobility affect p2p communications and they are unpredictable (CH₁), it is important to run test cases under various mobility patterns across all nodes in the network. Emulated mobility frameworks such as CORE [1] and EMANE [15] fall short as the connection latency and wireless transmission are specific to device model and radio. This is not a challenge in traditional mobile application development as the application functionality is confined within a single connected device.

D. Trust in information

As user devices with MNA advertise on unsecured wireless protocols, it may be possible for a malicious attacker to listen for such advertisements and reverse-engineer the message formats and protocols used. Then, the attackers may generate fake messages and advertise them, e.g., a fake tornado alert. Given a lack of trusted information in such scenarios (CH₃), misinformation campaigns can have disastrous consequences. In open decentralized systems, such false messages cannot be distinguished from the real ones, and MNA will end up propagating them to as many devices as possible, “poisoning” the network. In general, veracity of such information cannot be independently verified in open decentralized systems and previous works on peer-to-peer networks do not address this challenge.

III. C2P2P SYSTEM ARCHITECTURE




Fig. 1. c2p2p system architecture

To address (or bypass) the challenges identified above, this paper proposes c2p2p – a hybrid architecture that leverages a central service as the sole source of trusted information as depicted in Figure 1. An application-specific central service, e.g., weather.com backend, originates new information as a push message and assigns it a unique identifier. Next, the central service produces a digital signature using its private key and appends it to the message payload. Finally, header parameters specifying TTL (duration after which the message expires), peer identifier of the central service, destination peer identifier list (for uni-cast) or a special identifier (for broadcast), and the time of origin are added to the message and the message is sent to a push notification service for distribution.

Mobile applications integrate MNA as an SDK including the public key of the corresponding central service. Peers having Internet connectivity receive the push messages and verify the digital signature to protect against spread of misinformation. If the message has not expired and is destined to other peers, then the receiving peers start forwarding the message to other peers. Depending on the protocol being used, such forwarding happens as a uni-cast or broadcast.

Since new information can arrive at any time or new peers may come into proximity, and message transmissions are not reliable, all peers store unexpired and verified messages and continue to forward them to proximal peers repeatedly. This implies that peers may receive the same message more than once but since the messages are globally unique, duplicates can be identified and ignored. A peer sending a message adds its own peer identifier to a list of forwarders appended to the message header. A receiving peer thus knows all the peers that already have the message and stops repeating the message to those peers, controlling the flooding. As there are no control messages or any other overhead for routing, we call this *zero-overhead* routing.

The following assesses the c2p2p architecture relative to the challenges identified above along with known limitations. Further, to realize this architecture on mobile devices, we evaluated a variety of protocols on both Android and iOS. As described later, WiFi DNS on Android, Bluetooth Nearby on Android, and Bluetooth low energy (BTLE) on both Android and iOS are the only protocols that were found effective. Actual techniques for discovery, advertisement, and connection vary across these protocols and are described next.

A. Operating system restrictions

MNA works around the background activity and user permission restrictions via innovative techniques, without resorting to hacks that may violate user security or App-store guidelines. Tables I and II summarize the various protocols we have experimented with using WiFi and Bluetooth interfaces, along with key findings as described here. The protocols actually deployed are in bold.

TABLE I
ANDROID PROTOCOLS AND KEY FINDINGS

	Protocol	Throughput Avg. kbps	Battery per hour	Range LoS ft	iOS	Issues
WiFi	DNS	0.2	2%	600	No	Needs DNS Permission
	WiDi	2000	3%	600	No	
	Hotspot	2000	5%	600	No	
BT	Classic	50	2%	600	No	
	Nearby	50	2%	600	No	
	BTLE	50	2%	600	Yes	

TABLE II
IOS PROTOCOLS AND KEY FINDINGS

	Protocol	Throughput Avg. kbps	Battery per hour	Range LoS ft	Android	Issues
WiFi	Bonjour	0.2	5%	200	No	Battery Battery
	MPC	2000	5%	200	No	
	BTLE	50	1%	800	Yes	

Line of sight range in feet was measured by manually placing devices at increasing distances until they fail to communicate. Battery consumption per hour was measured as a difference between baseline battery consumption and when MNA was left running for 12 hours. Reported numbers are an average percentage consumption across about 50 devices on each platform.

BTLE is the only cross-platform protocol, although there are limitations with operating system versions as shown in Table III. Android is denoted as A, 18 and 21 are correspond to Android versions 4.3 and 5.0 respectively, and C and P denote the roles of Central and Peripheral. Key finding is that Android 4.3 devices acting as Peripheral cannot communicate either as senders or receivers.

TABLE III
BTLE INTEROPERABILITY BETWEEN ANDROID AND IOS DEVICES

Sender	Receiver					
	A.18-C	A.18-P	A.21-C	A.21-P	iOS-C	iOS-P
A.18-C		No		Yes		Yes
A.18-P	No		No		No	
A.21-C		No		Yes		Yes
A.21-P	Yes		Yes		Yes	
iOS-C		No		Yes		Yes
iOS-P	Yes		Yes		Yes	

WiFi Direct (WiDi) is a clever combination of Hotspot with WiFi DNS wherein peers randomly choose to be a Hotspot or a client, discover WiFi credentials over WiFi DNS discovery and make WiFi connection to Hotspot without needing user permission. However, both Hotspot and WiDi suffer from higher battery consumption than DNS and hence were not deployed. All Android protocols leverage foreground services to continue discovery and advertisement operations in the background indefinitely. Algorithms for WiFi DNS and BTLE are described in detail in Section IV.

B. Power constraints

Due to indefinite foreground services and continuous discovery and advertisement, this is primarily an issue on Android. Our approach here is two-pronged. Firstly, via extensive

experiments on a large number of devices, we fine-tune the algorithms governing the intervals at which discovery and advertisements occur. In a nutshell, receiving new information during a period causes MNA to be more aggressive in discovery and advertisement. Similarly, lack of new information for a period makes the device less aggressive. Secondly, we allow the central service to broadcast “wake up” messages ahead of an anticipated severe weather event. When devices receive such messages, they schedule themselves to remain aggressive during the specified window of time. Outside of this window, the device can afford to have long sleep cycles and conserve power. With these techniques, our testing shows less than 2% battery consumption per hour on most device models.

C. Testing p2p networks

We developed novel test automation tools and processes that control multiple devices from a single test station and follow prescribed steps to generate, send, and receives messages to play out a test scenario. The framework allows automated analysis of the observations to determine the test result. This capability was instrumental in uncovering bugs at a fast pace with tens of physical devices employed for automated testing. Further, the automation framework is general and can be expressly applied to test other apps in this fashion.

Figure 2 shows the overall test automation architecture. A tester specifies high-level test cases and expected results in a platform and device agnostic fashion. Mobile devices run a UI proxy for executing platform and device-specific UI command while a central laptop terminal acts as the overall orchestrator and translates high-level test cases and distributes them to devices. Since the orchestrator needs a wired access to the devices and the devices cannot be programmed to move (unlike drones and bots), testing of range and mobility patterns needs to be done manually, which is a known limitation.

A key innovation here is the platform-independent language for specifying high-level test cases and a small set of primitive commands it translates to. We describe here the list of primitive commands and a screen recording across seven Android devices demonstrates this powerful capability [9].

- `syncAll` Wait for all devices to reach this point
- `object.find` Find a UI object by name or position
- `object.click` Click a UI object
- `sleep` Wait for the specified duration
- `setupSystemAlert` Respond to a system popup
- `text.set` Highlight a text input box
- `text.write` Write text to the text box
- `swipe` Swipe to the left/right or up/down

D. Trust in information

In c2p2p, since the origin of information is always the central service, digital signatures can be attached to all messages originating from the service. The mobile application is distributed with the corresponding public key so that digital signatures from the service can be verified. If a message fails such a verification, it is discarded and not forwarded any further.

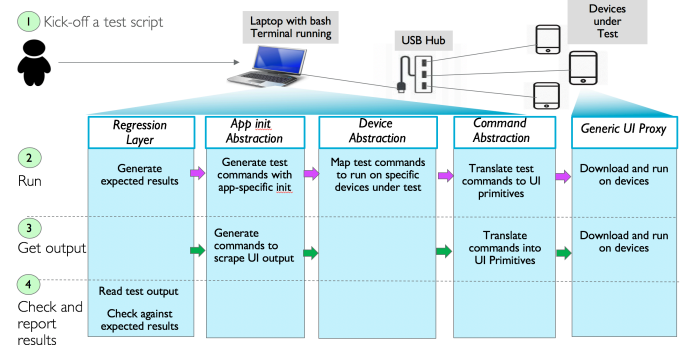


Fig. 2. Automated testing system for MNA

Peers cannot originate messages however as then a malicious peer can join the network, exchange keys with other peers, start generating, signing, and spreading fake information. This has been a longstanding problem for open decentralized systems. Requiring all peers to register their identity with a certificate authority and reporting the originators of fake information may discourage malicious behavior. Clearly, this is even harder to solve than the spread of fake information over the Internet because ground truth is hard to find and malicious peers can report genuine sources of information as well.

IV. ALGORITHMS

A. WiFi DNS

DNS service discovery is a widely supported protocol [4]. On Android, WiFi DNS improvises on this standard such that the exchange of information happens without making network connections at all. This is achieved by splitting the message payloads into small enough chunks and stuffing them into service advertisements as txt records and broadcast over multiple advertisements in a quick succession. Since there are no connections being made, the messages can be exchanged without prompting the user to grant permissions for each connection. Since no peer connections are needed and advertisements and discovery are broadcast, WiFi DNS is a broadcast protocol. On iOS, Bonjour protocol is built on DNS service discovery but is power-hungry and does not inter-operate with WiFi DNS on Android. Algorithm 1 outlines the peer process for WiFi DNS. Power conservation via time intervals is achieved by manipulating T_{ad} and T_{disc} based on traffic.

Messages from multiple applications are forwarded on the same network, but a peer can validate signatures only for the messages targeting an app it hosts. Peers simply forward messages not targeting one of their apps. Key benefit of this design is that the density of peers is available to all apps, making c2p2p scenarios viable even when an individual app does not have a sufficient density of users. However, this also allows a malicious jammer to generate traffic for a fake app and force the peers to forward it (as they do not match any of the apps) until the messages expire. Although the fake information will never be delivered to the apps, this issue

can affect MNA throughput for legitimate apps. A possible solution would be to maintain a central registry of a_{id} and distribute it to all peers across apps.

Algorithm 1: WiFi DNS peer algorithm

```

Input: Peer id  $p_{id}$ 
Input: App id  $a_{id}$ , max advert/discovery intervals
            $max_{ad}, max_{disc}$ 
Input: min advert/discovery intervals  $min_{ad}, min_{disc}$ 
 $S_{name} \leftarrow \text{'MNA'}$ 
 $M \leftarrow \emptyset$ 
 $T_{ad} \leftarrow min_{ad}$ 
 $T_{disc} \leftarrow min_{disc}$ 
 $New \leftarrow 0$  foreach timer expiry  $T_{ad}$  do
     $M \leftarrow M - \text{expired messages}$ 
     $M_{sort} \leftarrow M$  sorted by least advertised first
     $B \leftarrow M_{sort}$  chunked by MTU size
    dnsAdvertise ( $S_{name}, p_{id}, B$ )
    if  $New = 0$  then
         $T_{ad} \leftarrow \min(T_{ad} \times 2, max_{ad})$ 
    end
end
foreach timer expiry  $T_{disc}$  do
    dnsDiscover ( $S_{name}, p_{id}$ )
    if  $New = 0$  then
         $T_{disc} \leftarrow \min(T_{disc} \times 2, max_{disc})$ 
    end
     $New \leftarrow 0$ 
end
foreach discovered service callback  $S_{disc}$  do
     $B \leftarrow S_{disc}.B$ 
     $M_{part} \leftarrow msg(B)$ 
     $M \leftarrow M + M_{part}$ 
    if  $M_{part}$  is not a duplicate then
         $T_{ad} \leftarrow \max(\frac{T_{ad}}{2}, min_{ad})$ 
         $T_{disc} \leftarrow \max(\frac{T_{disc}}{2}, min_{disc})$ 
         $New \leftarrow 1$ 
    end
    if  $M_{part}$  was the last chunk of  $M$  then
        if  $M$  is unexpired then
            if  $M.app = a_{id} \wedge M$  signature is valid then
                Notify app with  $M$ 
            end
             $M \leftarrow M \cup M$ 
        end
    end
end
end
/* For testing only
foreach app request to send  $M$  do
    if  $M$  has not expired then
         $M \leftarrow M \cup M$ 
    end
end

```

B. BTLE

The only protocol that stays reliably active in background without violating iOS developer guidelines is BTLE. Unlike Android, iOS applications do not need to request discovery and advertisements continually. Once a BTLE Peripheral (role that has information) is advertised and a Central (role that receives information) requests discovery, the app is allowed to go to sleep. When a matching central or a peripheral device is found, iOS wakes up the application app to handle such an event, even when the application is in the background. MNA builds on this and leverages asynchronous dispatch queues to turn each peer into a Central as well as a Peripheral simultaneously. This architecture allows bidirectional data transfers without concurrency issues.

Algorithm 2 outlines the peer process for BTLE. The biggest difference with WiFi DNS is that discovery and advertisements happen only once per app launch. Secondly, the entire peer process is event-driven, so the app can sleep when there are no events to process. The list of forwarders for each message $M.fwdList$ is key in controlling the duplicates and in turn network utilization without overhead messages. In practice, limiting number of Peripherals a Central connects to enables reliable communications, specified by max_p .

V. EXPERIMENTAL RESULTS

Although MNA is deployed to about 10 million users, collecting a significant amount of data runs into privacy and scale issues. Also, given the challenges associated with tests involving a large number of physical devices along with human mobility patterns, we employ a three-pronged approach to evaluation: (1) Limited metrics from the MNA user population, (2) recruit human volunteers (colleagues at IBM) in a pilot to have the app on their personal devices and run tests in the background, and (3) use the automated test environment to run controlled experiments. All experiments mentioning WiFi are on the WiFi DNS channel and BT is Bluetooth classic.

A. MNA limited experiment



Fig. 3. Messages received over mesh in one week

As Figure 3 shows, measures of number of messages received first via MNA and the number of unique users was

Algorithm 2: BTLE peer algorithm

Input: Peer id p_{id} , App id a_{id} , Max peripherals max_p
 $S_{name} \leftarrow \text{'MNA'}$
 $\mathbb{M} \leftarrow \emptyset$
 $C \leftarrow \text{new Central() in new dispatch queue}$
 $P \leftarrow \text{new Peripheral() in new dispatch queue}$
 $T \leftarrow \text{new data sender dispatch queue}$
 $C.\text{bleDiscovery}(p_{id}, S_{name})$
 $P.\text{bleAdvert}(p_{id}, S_{name})$
 $P_{con} \leftarrow 0$
foreach *Peripheral peer found event* **do**
 if $P_{con} < max_p$ **then**
 $C.\text{connect}(peer)$
 $P_{con} \leftarrow P_{con} + 1$
 end
end
foreach *Peripheral peer connected event* **do**
 foreach $M \in \mathbb{M}$ **do**
 if $peer \notin M.fwdList$ **then**
 $T.\text{send}(M, peer)$
 end
 end
end
foreach *Peripheral peer disconnected event* **do**
 $P_{con} \leftarrow P_{con} - 1$
end
foreach *Central peer connected event* **do**
 foreach $M \in \mathbb{M}$ **do**
 if $peer \notin M.fwdList$ **then**
 $T.\text{send}(M, peer)$
 end
 end
end
foreach M *received event from* $T.\text{receive}()$ **do**
 if M *is not a duplicate* \wedge M *is unexpired* **then**
 if $M.app = a_{id} \wedge M$ *signature is valid* **then**
 Notify app with M
 end
 $M.fwdList \leftarrow M.fwdList \cup p_{id}$
 $\mathbb{M} \leftarrow \mathbb{M} \cup M$
 $T.\text{send}(M, \text{all connected peers})$
 end
end
 /* For testing only */
foreach *app request to send* M **do**
 if M *has not expired* **then**
 $M.fwdList \leftarrow p_{id}$
 $\mathbb{M} \leftarrow \mathbb{M} \cup M$
 $T.\text{send}(M, \text{all connected peers})$
 end
end

collected during a week in June 2018. Almost no MNA messages were received with a roll-out to 10% of the users.

This makes sense given that p2p needs critical density of users to propagate messages over large distances. With 100% of the users participating, 6 million+ weather alerts were received first via MNA by 1.9 million+ users within the collection period. Even when the users are connected to the Internet, they may receive the alert via MNA before the push notification arrives via the Internet. Given that such a large population of users received it via MNA first implies that there is value in the c2p2p architecture even when Internet connectivity is available.

B. Pilot test results

About 25 colleagues participated in a pilot on a site across two office buildings and about a thousand people in total. Clearly, the user density in this case was low. The app installed on their personal devices sent a broadcast message of varying sizes every five minutes and statistics on MNA were collected.

Overall, about 13% of the messages were received at least once (depends on density of users in the building) while users actively used the devices for 5% of the time. Figure 4 shows number of messages received and how the single-hop latency varies between WiFi DNS and BT Classic and across increasing packet sizes. Except for the 1000 byte+ packets, WiFi DNS achieved lower end-to-end latencies than BT Classic, perhaps due to its broadcast model and wide support across devices. This also explains why more messages were delivered over WiFi DNS than BT Classic. Packet sizes did not affect single-hop latencies on p2p DNS, perhaps because most of the messages fit within a single advertised chunk of up to 750 bytes.




Fig. 4. Latency as a function of packet sizes

C. Automated test results

We turn to the automated testing with up to 10 physical devices to carry out some controlled testing and a deeper understanding of how the various MNA protocols behave. Figure 5 shows end-to-end latency results in a test involving three Android devices, each sending a variety of message and payload size configurations. Latency is shown on a logarithmic scale y-axis. Total load on network increases from left to right on x-axis while the number of messages is small for the first




Fig. 5. End-to-end latency, WiFi DNS vs. BT Classic

and third configurations and large for the second and fourth configuration.

As expected, WiFi DNS latency increases significantly with network load given its very low throughput. As BT Classic cannot broadcast, it spends significant time making connections with peers. Hence, WiFi DNS beats BT Classic in the 30x67 configuration involving a large number of small messages. However, when the message sizes are increased as in 42x117 configuration, BT Classic is almost 10x faster than WiFi DNS due to superior throughput. Figure 6 confirms this difference in network topology by measuring average hop counts. WiFi DNS always needs fewer hops to deliver messages than BT Classic.




Fig. 6. Hop counts, WiFi DNS vs. BT Classic

Next, we compare WiFi DNS and BT Nearby on Android with BTLE on iOS in terms of end-to-end latency as the payload size is increased as shown in Figure 7. The results were collected over three repeats of the experiment involving up to 10 android devices and 6 iOS devices. BTLE is almost 100x faster than WiFi DNS and BT Nearby. For Wifi DNS, as payload size approaches the max advertisement chunk size of 750 bytes, the end-to-end latency rises by almost 10x (100 bytes versus 720 bytes). Behavior of BT Nearby taking time to form connections but having a good throughput is evident as the end-to-end latency does not change even when payload

sizes go up to 10KB. BTLE on the other hand gets 10x slower with 10KB payloads as compared to 1KB.




Fig. 7. End-to-end latency, WiFi DNS vs. BT Nearby vs. BTLE

As described above, MNA employs zero-overhead routing with the goal of maximizing coverage and minimizing latency across peers while tolerating topology changes. As a result, messages may be delivered multiple times to the same peer as no control messages are sent to keep track of the peer state. Figure 8 shows the cost of such an approach in terms of average number of duplicate messages received by peers across WiFi DNS, BT Nearby, and BTLE (iOS).

For small payload sizes, WiFi DNS ends up duplicating a large number of messages as the broadcasts are seen by many peers and repeated advertisements must go on to reach newly proximal peers. The duplicates drop significantly with payloads approaching the advertisement chunk size because the fraction of transmissions that are successful also drops with large payloads (multiple chunks need to be delivered per message). The forward list in BT Nearby and BTLE help control the duplicates significantly.




Fig. 8. Duplicity due to zero-overhead routing

Lastly, Figure 9 shows the network topology underlying WiFi DNS, BT Nearby, and BTLE. As shown earlier in Figure 6, WiFi DNS delivers most of the message with a single hop due to a broadcast protocol whereas BT Nearby and BTLE need as many as two hops on average.




Fig. 9. Average hop count

VI. RELATED WORK

Apart from the state-of-the-art cited earlier, there have been other notable attempts at practical large-scale deployments, some of which will be described here. Before mid-2000s, most of the research on MANETs was based on Department of Defense requirements, until commodity multi-hop ad hoc networks began to be considered [2]. Around that time, academia and industry started publishing and implementing commodity mesh networks in university and industrial environments. Some notable projects from this era are: the community-based multi-hop wireless network by Bahl et al. at Microsoft Research [14], the Broadband and Wireless Network (BWN) project at Georgia Tech [7], the testbed created at UCLA Network Research Lab for Wireless Mesh Networks and the evaluation of its performance of real-time traffic [3], the wireless mesh network testbed at Shanghai Jiao Tong University [20], and a survey paper about Wireless Mesh Networks and its challenges at the transport layer [19].

Despite these efforts, it took until 2014 before wireless mesh networking was used commercially to enable smartphones to connect via Bluetooth and WiFi in a popular application called FireChat [16]. The success of FireChat, partially due to the news coverage of its use in political situations in which governments restricted access to the Internet, has led to many alternatives in the past few years. An up-to-date list of such applications is available here⁴. With the recent introduction of p2p support at the Operating System level by Apple and Google for smartphones, we believe there will be a significant increase in both academic and industrial implementations of large scale mesh networks in the upcoming years.

VII. CONCLUSIONS AND FUTURE WORK

We reported on MNA, a large-scale real-world hybrid ad hoc network for Android and iOS platforms. Leveraging a central service as the single source of trusted information with a novel c2p2p architecture enables a class of applications to take advantage of p2p communications over an increasing density of mobile devices. Operating system restrictions are addressed or bypassed with simple and creative algorithms without violating user privacy and security or app developer

guidelines. Extensive experiments with real users, pilot users, and automation show that MNA enables low throughput connectivity on modern devices with low battery consumption on both Android and iOS platforms.

Although MNA represents a significant step forward in the right direction, there are significant limitations to be overcome. First, many applications require peers to originate information, e.g., asking for help during a natural disaster. Hence, the issue of trust in open decentralized systems needs to be investigated. Second, although MNA delivers almost all messages sent to all peers, it cannot guarantee a 100% delivery rate. In many applications, such guarantees are needed. These are challenging problems for further research.

ACKNOWLEDGMENT

This research was sponsored by the U.S. Army Research Laboratory and the U.K. Ministry of Defence under Agreement Number W911NF-16-3-0001. The views and conclusions contained in this document are those of the authors and should not be interpreted as representing the official policies, either expressed or implied, of the U.S. Army Research Laboratory, the U.S. Government, the U.K. Ministry of Defence or the U.K. Government. The U.S. and U.K. Governments are Authorized to reproduce and distribute reprints for Government purposes notwithstanding any copyright notation hereon.

REFERENCES

- [1] J. Ahrenholz, C. Danilov, T. R. Henderson, and J. H. Kim. Core: A real-time network emulator. In *MILCOM 2008 - 2008 IEEE Military Communications Conference*, pages 1–7, Nov 2008.
- [2] R. Bruno, M. Conti, and E. Gregori. Mesh networks: commodity multihop ad hoc networks. *IEEE Communications Magazine*, 43(3):123–131, March 2005.
- [3] V. Chavoutier, D. Maniezzo, C. E. Palazzi, and M. Gerla. Multimedia over wireless mesh networks : Results from a real testbed evaluation. In *Sixth Annual Mediterranean Ad Hoc Networking Workshop*, pages 56–62, June 2007.
- [4] S. Cheshire and R. Krochmal. Dns-based service discovery. *RFC 6763*, February 2013.
- [5] T. Clausen and P. Jacquet. Optimized link state routing protocol (olsr). *RFC 3626*, October 2003.
- [6] P. Gardner-Stephen and S. Palaniswamy. Serval mesh software-wifi multi model management. In *Proceedings of the 1st International Conference on Wireless Technologies for Humanitarian Relief*, ACWR '11, pages 71–77, New York, NY, USA, 2011. ACM.
- [7] Georgia Tech. <http://bwn.ece.gatech.edu/mesh/>.
- [8] IBM. Mesh Network Alerts (MNA). Available at <https://weather.com/apps/ibm/meshnetworkalerts>.
- [9] IBM. Mesh network alerts test automation demo. Available at <https://bit.ly/2Jj2wb>.
- [10] W. Kiess and M. Mauve. A survey on real-world implementations of mobile ad-hoc networks. *Ad Hoc Networks*, 5(3):324 – 339, 2007.
- [11] J. Loo, J. L. Mauri, and J. H. Ortiz. *Mobile Ad Hoc Networks: Current Status and Future Trends*. CRC Press, Inc., Boca Raton, FL, USA, 1st edition, 2011.
- [12] S. Marti, T. J. Giuli, K. Lai, and M. Baker. Mitigating routing misbehavior in mobile ad hoc networks. In *Proceedings of the 6th Annual International Conference on Mobile Computing and Networking*, MobiCom '00, pages 255–265, New York, NY, USA, 2000. ACM.
- [13] M. Mauve, J. Widmer, and H. Hartenstein. A survey on position-based routing in mobile ad hoc networks. *IEEE Network*, 15(6):30–39, Nov 2001.
- [14] Microsoft Inc. <https://www.microsoft.com/en-us/research/project/self-organizing-wireless-mesh-networks>.
- [15] Naval Research Laboratory. Extendable mobile ad-hoc network emulator (emane). Available at <http://cs.itd.nrl.navy.mil/work/emane/>.

- [16] Open Garden Inc. <https://www.opengarden.com/firechat/>.
- [17] C. Perkins, E. Belding-Royer, and S. Das. Ad hoc on-demand distance vector (aodv) routing. *RFC 3561*, July 2003.
- [18] C. E. Perkins. In *Ad Hoc Networking*, chapter Ad Hoc Networking: An Introduction, pages 1–28. Addison-Wesley Longman Publishing Co., Inc., Boston, MA, USA, 2001.
- [19] S. Saha, U. Kumar Acharjee, and M. Tahzib-Ul-Islam. A survey on wireless mesh network and its challenges at the transport layer. *International Journal of Computer Engineering and Technology (IJCET)*, 5:169–177, 08 2014.
- [20] Y. Wu, D. Sun, P. Yi, and J. Tang. Design and implementation of wireless mesh network testbed sjtu-mesh. In *2010 International Conference on Internet Technology and Applications*, pages 1–4, Aug 2010.
- [21] X. M. Zhang, Y. Zhang, F. Yan, and A. V. Vasilakos. Interference-based topology control algorithm for delay-constrained mobile ad hoc networks. *IEEE Transactions on Mobile Computing*, 14(4):742–754, April 2015.